

CPU Scheduling

Course Code: CSC 2209

Course Title: Operating Systems



Dept. of Computer Science
Faculty of Science and Technology

Lecturer No:	06	Week No:	06	Semester:	
Lecturer:	<i>Name & email</i>				

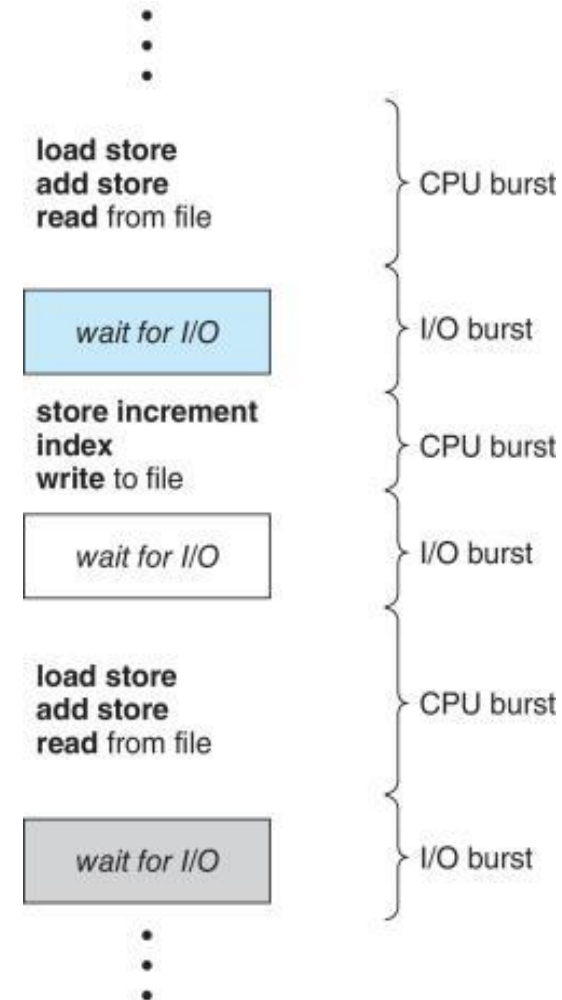
Lecture Outline



1. Basic Concepts
2. Scheduling Criteria
3. Scheduling Algorithms

Basic Concepts


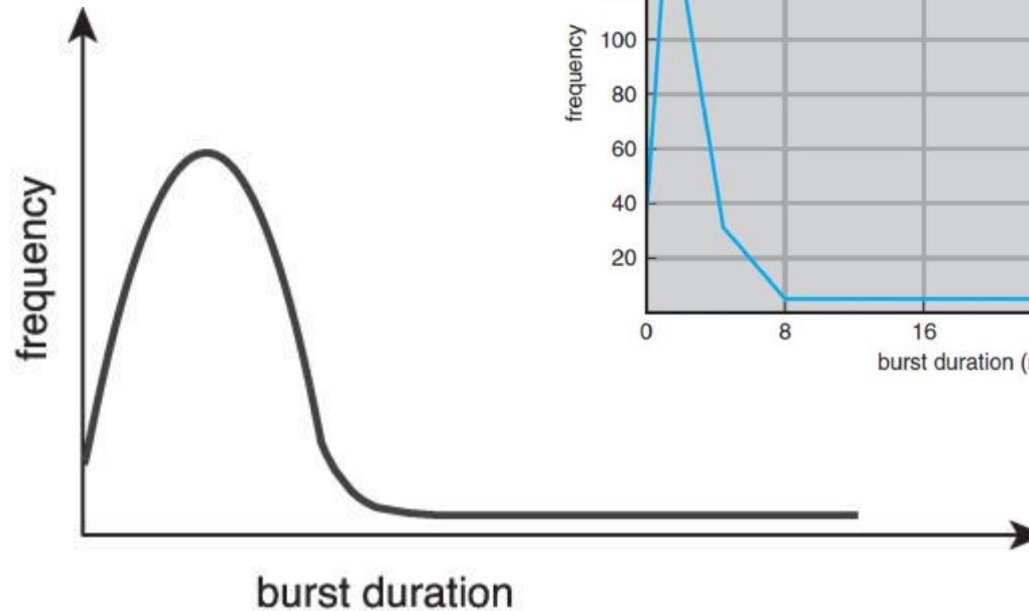
- ❑ Maximum CPU utilization obtained with multiprogramming
- ❑ CPU–I/O Burst Cycle – Process execution consists of a **cycle** of CPU execution and I/O wait
- ❑ **CPU burst** followed by **I/O burst**
- ❑ CPU burst distribution is of main concern



Histogram of CPU-burst Times


Large number of short bursts

Small number of longer bursts

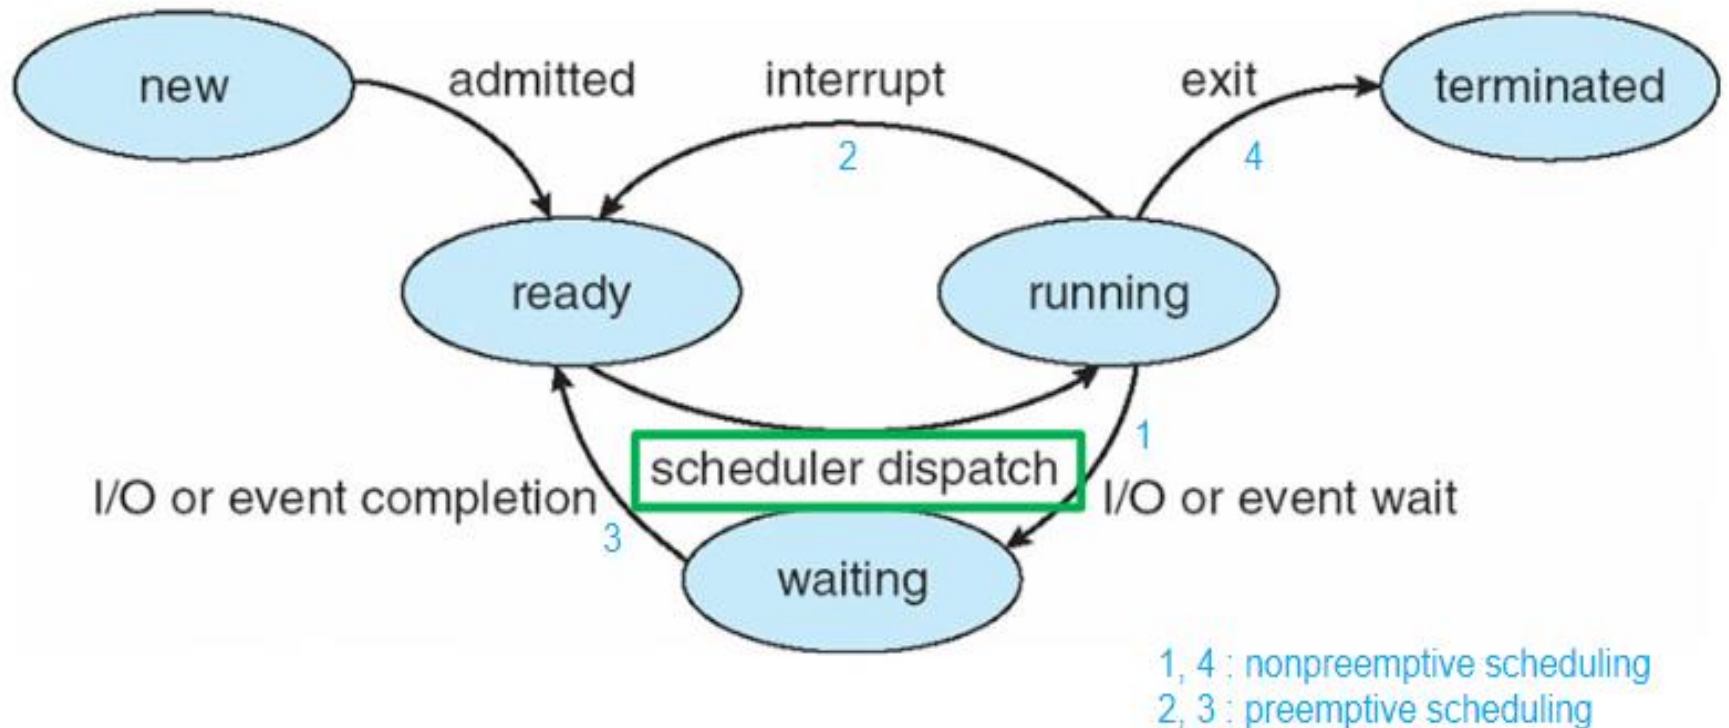


CPU Scheduler

- ❑ The **CPU scheduler** selects from among the processes in ready queue, and allocates the CPU to one of them
 - ❑ Queue may be ordered in various ways
- ❑ **CPU scheduling decisions may take place when a process:**
 1. Switches from running to waiting state
 2. Switches from running to ready state
 3. Switches from waiting to ready
 4. Terminates
- ❑ Scheduling under **1** and **4** is **nonpreemptive**
- ❑ All other scheduling is **preemptive**
 - ❑ Consider access to shared data
 - ❑ Consider preemption while in kernel mode
 - ❑ Consider interrupts occurring during crucial OS activities




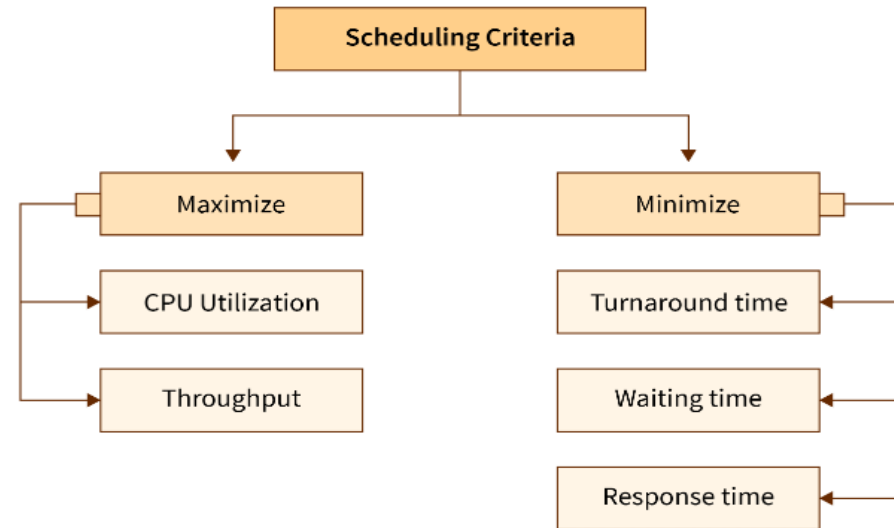
CPU Scheduler



- Many processes are in **ready** or **waiting** states in multiprogrammed operating systems.

Scheduling Criteria

- ❑ **CPU utilization** – keep the CPU as busy as possible
- ❑ **Throughput** – # of processes that complete their execution per time unit
- ❑ **Turnaround time (TAT)** – amount of time to execute a particular process
 - ❑ Completion time - Arrival time
- ❑ **Waiting time (WT)** – amount of time a process has been waiting in the ready queue
- ❑ **Response time (RT)** – amount of time it takes from when a request was submitted until the first response is produced, not output (for time-sharing environment)

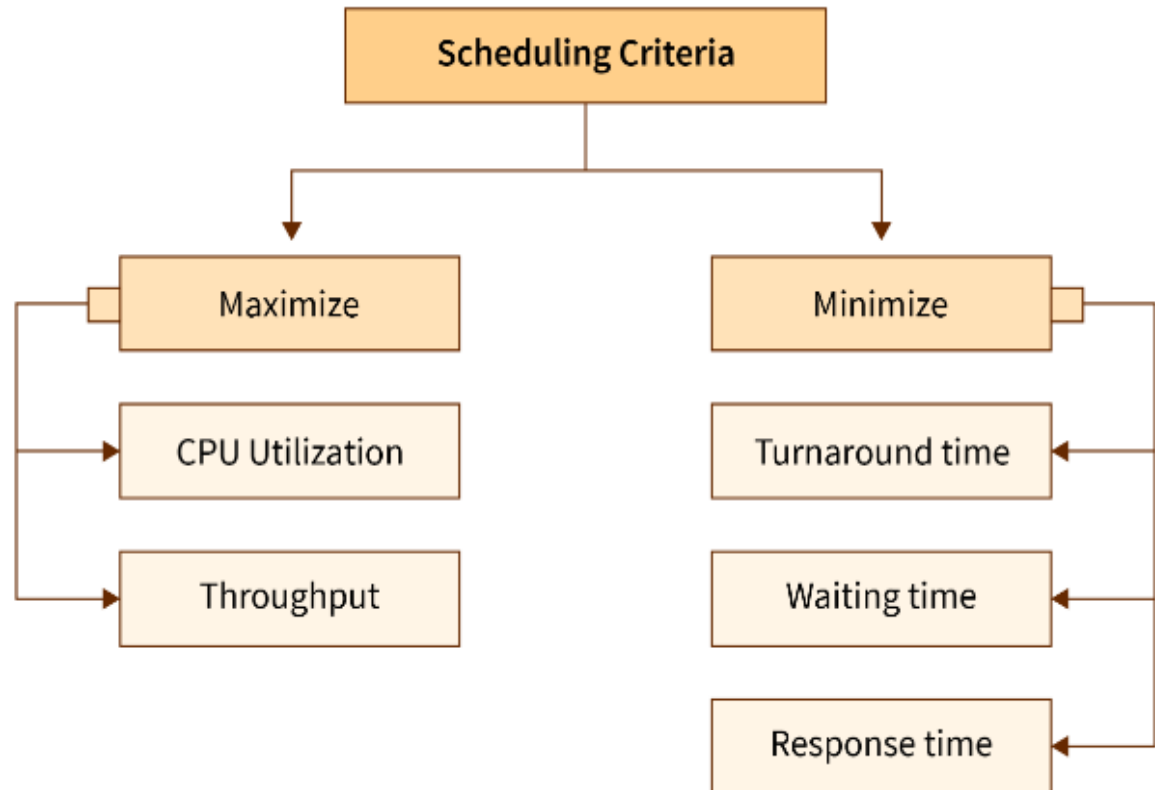


Turnaround time (TAT) = Completion time - Arrival time

Waiting time (WT) = Turnaround time - CPU burst time

Scheduling Algorithm Optimization Criteria

- ☐ Max CPU utilization
- ☐ Max throughput
- ☐ Min turnaround time
- ☐ Min waiting time
- ☐ Min response time



First- Come, First-Served (FCFS) Scheduling


<u>Process</u>	<u>Burst Time</u>
----------------	-------------------

P_1	24
-------	----

P_2	3
-------	---

P_3	3
-------	---

- ❑ Suppose that the **processes arrive in the order**: P_1, P_2, P_3
The Gantt Chart for the schedule is:



- ❑ Waiting time for $P_1 = 0$; $P_2 = 24$; $P_3 = 27$
- ❑ Average waiting time: $(0 + 24 + 27)/3 = 17$

FCFS Scheduling (cont'd)

- Suppose that the processes arrive in the order:

P_2, P_3, P_1

Process	Burst Time
P_1	24
P_2	3
P_3	3

- The Gantt chart for the schedule is:



- Waiting time for $P_1 = 6$; $P_2 = 0$; $P_3 = 3$
- Average waiting time: $(6 + 0 + 3)/3 = 3$
- Much better than previous case
- Convoy effect** - short process behind long process
 - Consider one CPU-bound and many I/O-bound processes


Shortest-Job-First (SJF) Scheduling

- ❑ Associate with each process the length of its next CPU burst
 - ❑ Use these lengths to schedule the process with the shortest time
- ❑ SJF is optimal – gives **minimum average waiting time for a given set of processes**
 - ❑ The difficulty is knowing the length of the next CPU request
 - ❑ Could ask the user

Example of SJF (Non preemptive)

<u>Process</u>	<u>Burst Time</u>
P_1	6
P_2	8
P_3	7
P_4	3

❑ SJF scheduling chart



❑ Average waiting time = $(3 + 16 + 9 + 0) / 4 = 7$

❑ Waiting time for $P_1 = 3$; $P_2 = 16$; $P_3 = 9$; $P_4 = 0$

Determining Length of Next CPU Burst

- ❑ Can only estimate the length – should be similar to the previous one
-> Then pick process with shortest predicted next CPU burst
- ❑ Can be done by using the length of previous CPU bursts, using **exponential averaging**

1. t_n = actual length of n^{th} CPU burst
2. τ_{n+1} = predicted value for the next CPU burst
3. $\alpha, 0 \leq \alpha \leq 1$ = weight of recent and past history
4. Define : $\tau_{n+1} = \alpha \cdot t_n + (1 - \alpha) \cdot \tau_n$

t_n : contains the most recent information

- ❑ Commonly, α set to $\frac{1}{2}$
- ❑ Preemptive version called **shortest-remaining-time-first**

τ_n : stores the past history

Determining Length of Next CPU Burst

Example:

Suppose we have the following sequence of recent CPU burst times for a process:

$$t_1 = 10 \text{ ms}, t_2 = 15 \text{ ms}, t_3 = 8 \text{ ms}, t_4 = 12 \text{ ms}$$

Assume $\alpha = 0.5$ (smoothing factor) and $\tau_0 = 10 \text{ ms}$ (initial estimate).

Calculation:

$$\tau_1 = 0.5 \cdot 10 + 0.5 \cdot 10 = 10 \text{ ms}$$

$$\tau_2 = 0.5 \cdot 15 + 0.5 \cdot 10 = 12.5 \text{ ms}$$

$$\tau_3 = 0.5 \cdot 8 + 0.5 \cdot 12.5 = 10.25 \text{ ms}$$



$$\tau_4 = 0.5 \cdot 12 + 0.5 \cdot 10.25 = 11.125 \text{ ms}$$

$$\tau_{n+1} = \alpha \cdot t_n + (1 - \alpha) \cdot \tau_n$$

Therefore, the prediction for the next CPU burst length τ_5 would be approximately 11.125 ms using this method.



Prediction of the Length of the Next CPU Burst

CPU burst (t_i)

"guess" (τ_i)

a set to $\frac{1}{2}$

$$8 = \frac{1}{2} \cdot 6 + \frac{1}{2} \cdot 10$$

$$6 = \frac{1}{2} \cdot 4 + \frac{1}{2} \cdot 8$$

Examples of Exponential Averaging

- $\alpha = 0$

- $\tau_{n+1} = \tau_n$

- Recent history does not count

- $\alpha = 1$

- $\tau_{n+1} = \alpha t_n$

- Only the actual last CPU burst counts

- If we expand the formula, we get:

$$\begin{aligned}\tau_{n+1} = & \alpha t_n + (1 - \alpha)\alpha t_{n-1} + \dots \\ & + (1 - \alpha)^j \alpha t_{n-j} + \dots \\ & + (1 - \alpha)^{n+1} \tau_0\end{aligned}$$


- Since both α and $(1 - \alpha)$ are less than or equal to 1, each successive term has less weight than its predecessor

Example of Shortest-remaining-time-first

- Now we add the concepts of **varying arrival times** and preemption to the analysis

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

- Preemptive SJF Gantt Chart*



- Average waiting time = $[(10-1-0)+(1-0-1)+(17-0-2)+(5-0-3)]/4 = 26/4 = 6.5$ ms
- Waiting time = **Total/Actual WT** – **previous total exec time** – **AT** (Formula)


Round Robin (RR)

- ❑ Each process gets a small unit of CPU time (**time quantum q**), usually 10-100 milliseconds. After this time has elapsed, the process is preempted and added to the end of the ready queue.
- ❑ If there are n processes in the ready queue and the time quantum is q , then each process gets $1/n$ of the CPU time in chunks of at most q time units at once. No process waits more than $(n-1)q$ time units.
- ❑ Timer interrupts every quantum to schedule next process
- ❑ Performance
 - ❑ q large \Rightarrow FIFO
 - ❑ q small $\Rightarrow q$ must be large with respect to context switch, otherwise overhead is too high

Example of RR with Time Quantum = 4


<u>Process</u>	<u>Burst Time</u>
P_1	24
P_2	3
P_3	3

❑ The Gantt chart is:




- ❑ Typically, **higher average turnaround than SJF**, but better *response*
- ❑ q should be large compared to context switch time
- ❑ q usually 10ms to 100ms, context switch < 10 usec

Time Quantum and Context Switch Time



Turnaround Time Varies With The Time Quantum



process	time
P_1	6
P_2	3
P_3	1
P_4	7

80% of CPU bursts
should be shorter than q

Priority Scheduling

- ❑ A priority number (integer) is associated with each process
- ❑ The CPU is allocated to the process with the highest priority (smallest integer \equiv highest priority)
 - ❑ Preemptive
 - ❑ Non-preemptive
- ❑ SJF is priority scheduling where priority is the inverse of predicted next CPU burst time
- ❑ Problem \equiv **Starvation** – low priority processes may never execute
- ❑ Solution \equiv **Aging** – as time progresses increase the priority of the process

Example of Priority Scheduling

<u>Process</u>	<u>Burst Time</u>	<u>Priority</u>
P_1	10	3
P_2	1	1
P_3	2	4
P_4	1	5
P_5	5	2

□ Priority scheduling Gantt Chart




□ Average waiting time = 8.2 msec

Priority Scheduling w/ Round-Robin

<u>Process</u>	<u>Burst Time</u>	<u>Priority</u>
P_1	4	3
P_2	5	2
P_3	8	2
P_4	7	1
P_5	3	3

- ❑ Run the process with the highest priority. Processes with the same priority run round-robin
- ❑ Gantt Chart with 2 ms time quantum





Books

- ❑ Operating Systems Concept
 - ❑ Written by Galvin and Silberschatz
 - ❑ Edition: 9th



References

- ❑ Operating Systems Concept
 - ❑ Written by Galvin and Silberschatz
 - ❑ Edition: 9th